TASK INTRODUCTION

You are working on a healthcare intelligence system that uses LLMs to assist with summarising patient data and generating structured recommendations for clinician review. You are given a short patient summary in plain text below. Your task is to design how an LLM should process this into a structured, clinician-readable output.

Patient Summary

58-year-old female with history of type 2 diabetes and hyperlipidemia. Latest labs (last 2–3 months):

- HbA1c: 7.8% (Jan), 8.4% (March)
- Fasting glucose: 6.5 mmol/L, 8.9 mmol/L (unclear which is most recent)
- LDL: 4.2 mmol/L

Medications:

- Metformin 500mg twice daily
- Atorvastatin 10mg (patient unsure if still taking)

Symptoms:

- Fatigue
- Occasional blurred vision

Notes:

- Patient reports improving diet recently
- Adherence unclear
- Some lab values recorded from different clinics

Your Task:

1. Prompt Design
   Write a prompt that instructs an LLM to generate a structured output for clinician review. Please account for factors a clinician would want to consider while structuring your prompt (Ex: no conflicting information)

2. Define the Output Structure
   Design a structured format that a doctor could quickly scan. You may use: JSON or clearly structured sections. Your structure should include (at minimum):
   - Key clinical summary
   - Trends (if applicable)
   - Risk flags
   - Uncertainties / data issues
   - Suggested follow-up areas

3. Produce a Sample Output
   Using your prompt (or manually), generate an example output.

4. Failure Modes, Mitigation & Prompt/Use Validation
   Identify what could go wrong (please be specific), and describe how you would detect and mitigate those risks. Examples of failure modes and usability risks include model hallucination or overconfidence. Please focus on human-in-the-loop designs.

Deliverables:

Please submit:

1. your prompt
2. your output structure
3. a sample output
4. a short write-up (~1 page or concise bullets)

What we are looking for:

- how you handle ambiguity and conflicting data
- how you structure LLM outputs
- your awareness of failure modes
- your ability to design safe, reliable systems

Submission Instructions:

Complete the above task and submit via email to eos.yixian@ora.group. Please indicate approximately how many hours you spent on this assignment. Successful candidates who clear the assessment round will be invited for up to a maximum of two short interviews with the hiring team.

Deadline:

30 April 2026